

Distributed Training Systems: Crash Pack

This is a single-volume condensation of Role Track 4 (Distributed Training Systems Engineer) and the base material it stands on. It holds the parts you need at your fingertips. The decision rules and the numbers and formulas worth memorizing. The traps that waste the most time. A rapid-fire self-test and a one-page cheat sheet.

How to use it. Read Section 1 once, slowly. It is the judgment layer: given a model and a cluster, what to pick and why. Use Section 2 as a reference while you work. It pairs every number with its units and its meaning, because a number without units is a rumor.

Read Section 3 before you debug anything; each trap names its symptom so you can match what you see. Work Section 4 with the answers covered. Print Section 5.

What this pack assumes. You have read the base volumes on distributed training, the two appendices on reliability and parallelism, and the Track 4 chapters. This pack does not re-teach. It compresses.

Scope note. Cluster work is specific to the hardware in front of you. The numbers here are the ones that survive across clusters: ratios, formulas, and order-of-magnitude thresholds. Where a value is site-specific, the text says so.

Section 1: Decision rules

Five decisions carry most of the role. Each rule below fits on an index card. The reasoning behind each one is in the track chapters; here is the operational form.

The vocabulary this section stands on:

Term	Meaning
rank	one GPU's training process: N ranks means N GPUs training together
training step	one forward pass (compute the model's outputs), one backward pass (compute gradients by flowing the error back through the layers), and one optimizer update
gradient	the direction each parameter must move to reduce the loss
optimizer state	what Adam keeps per parameter: two running numbers (momentum and variance)
microbatch	one slice of the global batch that flows through the pipeline stages
collective	one coordinated communication all ranks do together
all-reduce	every rank contributes its numbers and all end with the sum
all-gather	every rank collects everyone else's chunk
reduce-scatter	everyone gets a different chunk of the global sum
NVLink domain	the set of GPUs that talk over NVLink: usually the 8 GPUs inside one node

1.1 Parallelism strategy selection: memory first, then bandwidth

The order never changes: DDP, then FSDP, then tensor parallelism, then pipeline parallelism. You move down the list only when the current option fails a concrete test. Context and expert parallelism are not on this ladder; they answer different questions (sequence length, MoE routing) and layer on top.




Figure 4.2. MoE capacity arithmetic. Expected load is 128 tokens per expert. The buffer holds 160 (factor 1.25). When 384 tokens route to one hot expert, 224 drop. The balancer nudges the router away on later steps.




Figure 4.3. Ring attention versus Ulysses at long sequences. Ring attention overlaps block transfers with compute, so per-GPU memory stays flat. Ulysses briefly holds all tokens and OOMs at extreme lengths. Byte counts lose to overlap plus memory scaling.



How to read this diagram. Start at the top. Every diamond is a measurement, not an opinion. "Does it fit" is arithmetic on parameter bytes, optimizer bytes, and activation bytes. "Is comm under 15%" comes from the alpha-beta model (transfer time = fixed latency per step + bytes / bandwidth; section 2.2 works it fully) or a profile. "TP inside the NVLink domain" is a hardware fact about your cluster.

The last gate is the bubble formula from Section 2: $(p-1)/(m+p-1)$. The final answer always comes from running the candidate layouts and reading the trace. This ladder keeps you from trying the obviously wrong layouts first.

The parallelism decision ladder: add strategies only when forced





Figure 4.1. The decision ladder in one picture. Start at DDP. Move down only when the current option fails its test: memory fit, then comm fraction, then depth. Each box names the gate that forces the move.

How to read this diagram

1. Start at the top box (DDP): the default. It works when one GPU holds the whole model.
2. Drop to FSDP when memory runs out: sharding trades communication bytes for per-GPU memory.
3. Drop to TP when FSDP all-gather traffic passes about 15% of step time: TP moves that traffic onto fast NVLink.
4. Drop to PP only for very deep models that still do not fit: the pipeline bubble is the price.
5. What it is: a forced march, not a menu. Why it exists: each strategy fixes one bottleneck and adds one cost. What breaks: skipping a rung (TP across slow links) or climbing too early (FSDP on a model that fits).

The six axes in one table:

Axis	What is split	What you pay	When it wins
Data (DDP)	Batch across replicas	One gradient all-reduce per step	Model fits on one GPU
Sharded data (FSDP / ZeRO-3)	Batch plus params, grads, optimizer	All-gather plus reduce-scatter per step	Big model, limited per-GPU memory
Tensor (TP)	Individual layers across GPUs	All-reduce inside every layer, fwd and bwd	Fast intra-node links (NVLink)
Pipeline (PP)	Layers across stages	Idle bubble, point-to-point sends	Many layers, slow cross-node links
Context (CP)	Sequence across GPUs	Attention-block exchanges	Long sequences
Expert (EP)	MoE experts across GPUs	All-to-all dispatch and combine	MoE models








Figure 4.4. The ZeRO tradeoff. Memory per GPU falls from 120 GB (DDP) to 14 GB (ZeRO-3) for a 7B model. Communication per step rises from 1.75M to 2.6M bytes. Sharding is a trade, not a free win: fitting the model is the point.



FSDP is ZeRO-3 with on-demand materialization. Each GPU keeps one shard at rest; an all-gather rebuilds the full layer just for the forward pass, then frees it. Backward repeats the all-gather and reduce-scatters the gradients. The price is 1.5x the all-reduce traffic of DDP; prefetch hides it when each layer does enough compute.

The two traps this ladder exists to prevent. TP across slow inter-node links is almost always a loss. The per-layer collectives that cost 19 ms on NVLink cost 350 ms on InfiniBand, and no overlap hides that. PP with too few microbatches can idle nearly half the step: 8 stages with 8 microbatches is a 47% bubble.



How to read this diagram


1. Read the two dashed boxes as two nodes. Each holds 4 GPUs joined by NVLink.
2. Inside each node, the green label marks tensor parallelism: per-layer all-reduces at 19 ms on this fabric.
3. Between nodes, the red arrow marks data parallelism: one gradient all-reduce per step at 350 ms on the same bytes.
4. What it is: a placement rule. Why it exists: TP collectives are latency-sensitive and frequent, so they need the fast domain. What breaks: TP across nodes (19 ms becomes 350 ms), or DP inside the node (wastes NVLink on a tolerant op).

Key takeaways

- The ladder is DDP -> FSDP -> TP -> PP, gated by memory fit, comm fraction, NVLink domain, and bubble size.
- Every gate is a number. If you cannot compute it, you guess.

The parallelism strategy is set. Now the cost of talking between ranks: which collective algorithm, on which fabric.

1.2 Collective algorithm selection: ring for bandwidth, tree for latency



A 4-stage pipeline over 6 microbatches. Forward passes warm up diagonally, backward passes drain the same way, and the dashed cells are the bubble: stages waiting with nothing to do. Bubble fraction = $(p-1)/(m+p-1)$, here $3/9 = 33\%$. The pack's warning case: 8 stages with 8 microbatches idles 47% of the step. 1F1B trims the middle bubble by starting backward passes earlier.




Figure 4.6. The GPipe schedule. All four microbatches run forward first, then all run backward. Stages idle while the wave passes: a 50% bubble here. The idle area is arithmetic, not a tuning problem.



Figure 4.7. The 1F1B schedule. Each stage alternates one forward with one backward once the pipe fills. The bubble shrinks to about 25%. Same stages, same microbatches, less idle time.

How to read this diagram

1. Read each row as one pipeline stage, time running left to right. Green blocks are forward passes, blue blocks are backward passes.
2. In Figure 4.6, follow stage 0: four forwards, then a long wait, then four backwards. The wait is the bubble.
3. In Figure 4.7, follow stage 0: forward, backward, forward, backward. Backward starts as soon as the last forward of the next microbatch clears.
4. What it is: two ways to order the same work. Why it exists: 1F1B keeps memory lower and idle time smaller. What breaks: too few microbatches (the bubble formula $(p-1)/(m+p-1)$ dominates), or activation memory (1F1B holds fewer live activations than GPipe).

Every collective is three decisions: the operation (what the math needs), the algorithm (how ranks cooperate), and the protocol (how bytes cross one link). The rule NCCL (NVIDIA's GPU collective-communications library) itself uses:

**Phase 1: reduce-scatter.**

Each GPU sums one chunk around the ring. After 3 steps every GPU holds one fully reduced chunk.

**Phase 2: all-gather.**

Each GPU broadcasts its reduced chunk around the ring. Everyone ends with all four reduced chunks.

Bandwidth-optimal: each GPU sends $2 \cdot (N-1)/N$ * bytes. Tree wins only when latency, not bytes, dominates.

Ring all-reduce in two phases. Phase 1 (reduce-scatter): each GPU accumulates one chunk around the ring, so every GPU ends with one fully reduced chunk. **Phase 2 (all-gather):** each GPU broadcasts its chunk, so everyone ends with all chunks. **Bandwidth-optimal:** each GPU sends $2 \cdot (N-1)/N$ of the bytes. The tree algorithm wins only when latency, not bytes, dominates.

- **Small messages: tree + LL protocol.** The tree finishes in $2 \cdot \log_2(N)$ phases instead of $2(N-1)$ steps. LL moves 8 bytes per flag with microsecond latency via CPU polling.
- **Large messages: ring + Simple protocol.** The ring keeps every link busy in every step; each rank moves $2(N-1)/N \cdot M$ bytes, about $2M$ at large N . Simple is bulk DMA at full bandwidth.
- **In-network reduction (CollNet / NVLS) only where the hardware allows.** SHARP-capable InfiniBand switches (SHARP: the switch itself sums the data, so ranks move less traffic), or NVSwitch inside the node. Validate with it off first: misconfigured in-network reduction is a source of silent wrongness.

The crossover sits in the low megabytes on IB-class fabrics. On 64 ranks with 40 GB/s effective bandwidth and 5 us per step, the tree wins below about 2.3 MB. The ring wins above. The calculator in Section 2 recomputes this for your fabric.

When to override the tuner. Almost never, and only with a measurement. Two legitimate cases. Force Tree/LL to work around a sick ring path on your fabric. Force Ring/Simple when you know the message mix better than the tuner does. If forcing the tree fixes a hang, the ring path on your fabric is suspect, and that is a hardware ticket, not a config win.

The protocol table, for reference:

Protocol	Chunk on the wire	Strength	Weakness
Simple	Large (bulk DMA)	Highest bandwidth	Fence and sync latency per chunk
LL (low latency)	8 bytes	Lowest latency, CPU-pollled	Low bandwidth; burns host CPU
LL128	128 bytes	Middle ground	Needs 128-byte atomic writes (not all fabrics)

The protocol is not a quality setting. LL is not "worse" than Simple. For a 4 KB all-reduce across 64 ranks, tree + LL finishes in about 60 us while ring + Simple takes about 630 us. Latency dominates small messages. Bandwidth dominates large ones. The alpha-beta model decides, not the name.

Parallelize the run and choose the collectives. Now the insurance: how often to checkpoint, by Daly not vibes.

1.3 Checkpoint interval: Daly, not vibes

Let C be the stall cost of one checkpoint in seconds and μ the mean time between failures in seconds. The interval that minimizes expected waste is:

$$\tau^* = \sqrt{2 \times C \times \mu}$$

Where: τ^* is the optimal checkpoint interval in seconds, C is the save stall in seconds, μ (μ) is the mean time between failures in seconds. It balances expected lost work (grows with the interval) against write overhead (shrinks with the interval).

The shape behind it: failures land at random points in the interval, so average lost work per failure is $\tau/2$. That term rises with the interval. Write overhead per unit time is C/τ , which falls with the interval. Their sum makes a U shape. The minimum is where they balance.

Worked reference (70B, Adam, BF16, 1,000 GPUs): $C = 200$ s of stall, $\mu = 50.7$ h = 182,520 s.

$$\tau^* = \sqrt{2 * 200 * 182,520} = 8,544 \text{ s} = 142 \text{ min}$$

About every 2.4 hours. Expected loss per failure is $\tau/2 = 71$ min of work plus restart time. Cut the stall to 30 s with async checkpointing and the optimum drops to 55 min. Faster writes justify more frequent checkpoints. That is the whole point of async checkpointing expressed as a number.

The script computes the U-curve below; it is not hand-drawn. It shows waste per hour at several intervals for the same C and mu. The script that computes it:

```
# Daly U-curve: expected waste per hour vs checkpoint interval.
# WHY: the curve is computed from waste(tau) = C/tau + tau/(2*mu),
# not hand-drawn, so the 142-min optimum in the text and the bars
# below cannot disagree.
# WHAT BREAKS: C and mu are estimates; the curve's bottom is shallow,
# so the policy (checkpoint roughly every 1-3 h) survives bad
# estimates. Do not over-tune the interval.

C = 200.0          # seconds of write stall per checkpoint
mu = 50.7 * 3600   # MTBF in seconds (1k-GPU run)

def waste(tau_s):
    # Write overhead per unit time plus expected lost work per failure.
    return C / tau_s + tau_s / (2 * mu)

print(" interval  waste/hr")
for minutes in (15, 30, 60, 142, 240, 480, 960):
    tau = minutes * 60
    w = waste(tau)
    bar = "#" * int(round(w * 250))
    tag = " <-- optimum" if minutes == 142 else ""
    print(f" {minutes:>4} min  {w*100:5.2f}% {bar}{tag}")
```

Its executed output, read as a table with a picture attached. The bottom is shallow: 60 min or 240 min costs only slightly more than the 142-min optimum. Do not over-tune; C and mu are estimates anyway.

interval	waste/hr	
15 min	22.47%	#####
30 min	11.60%	#####
60 min	6.54%	#####
142 min	4.68%	##### <-- optimum
240 min	5.33%	#####
480 min	8.58%	#####
960 min	16.13%	#####

How to read this figure. Each bar shows expected waste as a fraction of run time: write overhead plus expected lost work per failure. The left side is the write tax (checkpointing every 15 min spends 22% of the run writing). The right side is failure exposure (checkpointing every 16 h loses 16% to rework). The optimum sits where the two marginal costs balance.


The shallow bottom makes the policy forgiving: being near right beats being exactly right.

Retention rule. Keep the last 3 checkpoints plus periodic older ones (one per day). Never keep exactly one: a single corrupt write loses everything. Storage is cheap next to the GPU-hours the checkpoint protects.

Checkpoints exist. Now the guardian loop that turns a failure into a resume in under ten minutes: detect, classify, isolate, resume.

Daily checkpoint interval: the U-curve is computed, not drawn

waste per hour = $C/\tau + \tau/(2\mu)$; $C = 200$ s stall, $\mu = 50.7$ h (1,000 GPUs)




The Daly U-curve, computed from $\text{waste} = C/\tau + \tau/(2\mu)$ with a 200 s write stall and 50.7 h MTBF. Checkpointing every 15 minutes spends 22% of the run writing; every 16 hours loses 16% to rework. The optimum sits at 142 minutes, and the shallow bottom means 60 or 240 minutes costs only slightly more: be near right, not exactly right.

What to look at first: the green bar at 142 min, the lowest waste. The left side is the write tax, the right side is failure exposure. The dashed line marks the optimum. The one takeaway: async checkpointing cuts the stall to 30 s and moves the optimum to 55 min. That is the whole point of async checkpointing as a number.

1.4 Recovery strategy: the guardian loop

Failures are the normal operating condition at scale, not an exception. The recovery strategy is a control loop with a time budget: detect plus isolate plus resume stays under 10 minutes.



How to read this diagram. The loop starts with heartbeats: every rank reports step completions and the guardian watches the distribution. An anomaly triggers classification because the response differs by cause.

A dead node gets quarantined and the job re-queues without it. A hang gets a kill and a clean resume. A loss spike gets a rollback to before the spike began, and a human gets paged after the second rollback. The loop closes back to monitoring.

Detection thresholds, concretely:

- **Hang:** no rank completes a step for longer than 5x the p99 step time (p99: the 99th percentile; 99% of steps are faster). A fixed wall-clock guess breaks when the workload changes.
- **Straggler:** one rank's step time persistently over the fleet median. Above 1.1x investigate, above 1.3x act.
- **Loss spike:** loss above 1.5x the running median for several consecutive logs. Roll back automatically.

- **Node gone:** heartbeat lost. Quarantine at once; do not wait to see if it returns.

The rollback protocol for loss spikes, in order:

1. **Stop and checkpoint nothing.** Current weights are suspect.
2. **Roll back** to the last checkpoint saved before the spike started.
3. **Skip past the bad batches** using the saved data position.
4. **Resume with the learning rate cut 2-4x** for a few hundred steps, then restore the schedule.
5. **If it spikes again at the same data position, the data is the problem:** investigate offline.




Figure 4.10. The persist counter in action. The dashed line is the $k = 10$ MAD threshold. A 2-step blip crosses it but the streak counter stays under 5, so training continues. A 5-step sustained spike crosses the persist gate: the guardian rolls back to the last checkpoint before the spike.

How to read this diagram


1. Read the x-axis as training steps, the y-axis as loss. The dashed red line is the MAD threshold: $k = 10$ above the running median.
2. The blue band marks a 2-step spike: a noisy batch. The streak counter reaches 2, under the persist gate of 5, so nothing fires.
3. The red band marks a 5-step spike: poisoned momentum buffers. The streak counter reaches 5, the gate fires, and the guardian rolls back.
4. What it is: a two-threshold detector. Why it exists: rolling back on every blip wastes hours; never rolling back wastes the run. What breaks: k tuned on the wrong baseline (false positives after a data change), or a persist gate of 1 (hair-trigger rollbacks).

Isolation discipline. Quarantined nodes get re-qualified with the bring-up tests, abbreviated, before they return to the scheduling pool. Nodes that fail re-qualification get a hardware ticket. Node-level failure concentration is real: in one study, 3 of 63 nodes caused over half of all exclusions.

The run is parallel, checkpointed, and recoverable. Last decision: which number format the math runs in, and which peak your MFU quotes.

The guardian loop: detect, classify, isolate, resume in under 10 minutes

failures are the normal operating condition at scale, not an exception



Budget: detect + isolate + resume stays under 10 minutes. Bad nodes never rejoin silently: re-qualify with bring-up tests first.

The recovery control loop: heartbeats feed a monitor. The monitor classifies anomalies because the response differs by cause. It quarantines suspect nodes, and they never rejoin silently. The run resumes from the latest good checkpoint with a 10-step verification. Detect plus isolate plus resume stays under 10 minutes total.

What to look at first: the top-down flow from monitor to anomaly to classify. The four colored boxes are the four failure causes and their signatures. The green arrow on the right closes the loop back to monitoring. The one takeaway: classification comes before action, because a hang, a slow rank, a loss spike, and a dead node each need a different response.


1.5 Precision selection: BF16 by default, FP8 with a recipe

Precision is a three-way trade between memory, speed, and numerical safety. The rule:


1. **Default to BF16.** It has the same exponent range as FP32 (standard 32-bit floating point, the full-precision default). Overflow behavior matches what most recipes assume. No loss scaler needed for most layers.

2. **Use FP16 only when BF16 is unavailable.** Its narrower range overflows sooner. If a NaN (not a number: the value a computation produces when it overflows) appears right after a precision change, the change is the prime suspect. BF16 is the safer drop-in.
3. **Use FP8 when the hardware and recipe both support it.** H100, B200, and TPU Ironwood have FP8 matmul support. FP8 halves the bytes of BF16 (memory and bandwidth win), but it needs careful per-tensor scaling and some layers (norms, the normalization layers; softmax; the loss) usually stay in higher precision. Transformer Engine-style libraries automate the scaling; validate the recipe before trusting it at scale.

The MFU honesty rule. MFU means achieved FLOPs divided by peak FLOPS, and the peak must match the precision you actually used. Quoting FP8-achieved FLOPs against the FP8 peak is correct; quoting BF16 math against the FP8 peak halves the MFU silently. Always state the peak: "MFU 55% against the 989 TFLOPS BF16 dense peak."

Key takeaways

Tensor parallelism stays inside the NVLink domain; data parallelism crosses nodes. This is the 3D-parallel layout the decision ladder builds toward. Figure: Shoeybi et al., Megatron-LM (arXiv:1909.08053).



The first GEMM's weight A is split by columns (A_1, A_2); each GPU computes its half, applies GeLU, then multiplies by its row-split half of B . The f and g ops are the all-reduces that stitch the halves back together. This is why TP needs fast intra-node links: every layer pays two all-reduces. Figure: Shoeybi et al., arXiv:1909.08053.


Key takeaways

- Parallelism: DDP -> FSDP -> TP -> PP, each gate a measurement.
- Collectives: tree + LL below the crossover, ring + Simple above; override the tuner only with a measurement.
- Checkpointing: $\tau^* = \sqrt{2 \cdot C \cdot \mu}$, retention $K = 3$ plus daily.
- Recovery: detect, classify, isolate, resume, under 10 minutes total; rollback protocol for spikes.
- Precision: BF16 default; FP8 only with a validated recipe; always quote the peak your MFU uses.

Next: Section 2, the numbers that survive a cluster change: throughput math, collective costs, checkpoint intervals, and the TPU figures.

Precision: BF16 by default, FP8 only with a validated recipe

same exponent range as FP32 is why BF16 is the safer drop-in



MFU honesty rule: the peak must match the precision actually used. Quoting BF16 math against the FP8 peak halves the MFU silently. Always state the peak: "55% against the 989 TFLOPS BF16 dense peak."

Why BF16 is the default: it keeps the same 8-bit exponent range as FP32. Overflow behavior matches what most recipes assume. FP16 narrows the exponent and overflows sooner; FP8 halves the bytes but needs per-tensor scaling, with norms and softmax kept in higher precision.

What to look at first: the blue exponent segments, identical in width for FP32 and BF16. Arrows are not needed here: compare segment widths across rows. The one takeaway: match the MFU peak to the precision you actually used, and always state it.

Section 2: Must-memorize numbers and formulas

The rules say what to pick. Section 2 gives you the arithmetic to check the pick: through-put and MFU, collective costs, the Daly interval, failure math, and the TPU/XLA numbers. It answers the question Section 1 leaves open: what are the exact quantities behind "does it fit" and "will it finish"?

Every entry gives the value, the units, and what it means operationally. These are the numbers that survive a cluster change.

2.1 Throughput and efficiency

#	Quantity	Value	Units / meaning
1	6ND rule	$6 * N_{\text{params}} * D_{\text{tokens}}$	FLOPs to train N parameters on D tokens. 2N forward, 4N backward. The numerator of every MFU calculation and every training-cost estimate.
2	MFU	$(6 * N * \text{tok/s}) / (n_{\text{gpus}} * \text{peak})$	Dimensionless ratio. Achieved model FLOPs over hardware peak. Counts only useful math: comm, recompute, and bubbles lower it.
3	H100 BF16 dense peak	989 TFLOPS	TeraFLOPs per H100 SXM. The denominator for BF16 MFU. A100 is 312 BF16 TFLOPS for comparison.
4	Healthy MFU target	50 to 65%	Dense transformer on H100-class hardware. Below 40% means profile now.
5	70B reference point	9,423 tok/s = 50% MFU	Global throughput on 8x H100, BF16. A healthy 70B run should beat this.
6	Pipeline bubble (1F1B)	$(p-1) / (m + p-1)$	Fraction of the PP schedule lost to fill and drain. p stages, m microbatches.
7	Bubble worked values	8x32: 18%; 8x8: 47%; 6x18: 22%	Read as: stages x microbatches. PP without enough microbatches idles nearly half the step.

8	Goodput	Useful tok/s after all overhead	What the model actually receives. A 60% MFU run with constant restarts can deliver less than a 50% MFU run that never stops.
9	Straggler factor	max / median rank step time	Above 1.1x investigate, above 1.3x act. One 30%-slow rank in a 1,000-rank job is a 30% tax on everything.

$$\text{MFU} = (6 \times N \times \text{tok/s} / n_{\text{gpus}}) / \text{peak}$$

Where: N is the parameter count. tok/s is global throughput in tokens per second. n_{gpus} is the GPU count. Peak is the hardware ceiling at the precision used (989 TFLOPS BF16 per H100). Dimensionless; counts only useful math.

Worked: the 70B reference point: $(6 \times 70\text{e}9 \times 9,423 / 8) / 989\text{e}12 = 4.947\text{e}14 / 9.89\text{e}14 = 0.50$, so 50% MFU on $8 \times \text{H100}$, BF16. A healthy 70B run should beat this.

```

import math

# Core calculators: MFU, pipeline bubble, FSDP communication volume.
# WHY: these three numbers answer "is the run healthy", "does PP make
#   sense here", and "is the interconnect the bottleneck".
# WHAT BREAKS if changed: using the FP8 peak instead of the BF16 peak
#   silently shifts every MFU. TFLOPS vs FLOPS shifts it by 1e12.

H100_BF16_PEAK = 989.4e12 # dense BF16 TFLOPS per H100 SXM, vendor spec

def mfu(n_params, tokens_per_sec, n_gpus, peak=H100_BF16_PEAK):
    # Achieved FLOPs come from the 6ND rule: ~6 FLOPs per parameter per
    # token (2 forward + 4 backward). Divide by hardware peak.
    achieved = 6.0 * n_params * tokens_per_sec
    return achieved / (n_gpus * peak)

def pipeline_bubble_fraction(n_stages, n_microbatches):
    # 1F1B schedule: fill takes (p-1) slots, drain the same. Total
    # schedule is (m+p-1) slots, so the idle fraction is (p-1)/(m+p-1).
    p, m = n_stages, n_microbatches
    return (p - 1) / (m + p - 1)

def fsdp_comm_per_rank_gb(n_params, n_ranks, bytes_per_param=2):
    # FSDP pays the 2(N-1)/N factor twice per step: all-gather of params
    # (forward + backward) and reduce-scatter of gradients.
    model_gb = n_params * bytes_per_param / 1e9
    return 2 * (2.0 * (n_ranks - 1) / n_ranks) * model_gb

print("== MFU: 70B dense, 8x H100, BF16 ==")
N, GPUS = 70e9, 8
for tok_s in (4000, 9423, 12000):
    # 9423 tok/s is constructed to land exactly on 50%: the throughput
    # a healthy 70B run on 8 H100s should beat.
    print(f" {tok_s:>6} tok/s -> MFU {mfu(N, tok_s, GPUS):.1%}")

print("== pipeline bubble (1F1B) ==")
for stages, micros in ((8, 32), (4, 16), (8, 8)):
    b = pipeline_bubble_fraction(stages, micros)
    print(f" {stages} stages x {micros:>2} microbatches -> bubble {b:.1%}")

print("== FSDP comm per rank per step (70B, BF16) ==")
for ranks in (64, 512):
    gb = fsdp_comm_per_rank_gb(70e9, ranks)
    print(f" {ranks:>3} ranks -> {gb:6.1f} GB per rank per step")

```

The executed output of the script above is:

```

== MFU: 70B dense, 8x H100, BF16 ==
  4000 tok/s  ->  MFU 21.2%
  9423 tok/s  ->  MFU 50.0%
 12000 tok/s  ->  MFU 63.7%
== pipeline bubble (1F1B) ==
  8 stages x 32 microbatches -> bubble 17.9%
  4 stages x 16 microbatches -> bubble 15.8%
  8 stages x  8 microbatches -> bubble 46.7%
== FSDP comm per rank per step (70B, BF16) ==
  64 ranks ->  551.2 GB per rank per step
 512 ranks ->  558.9 GB per rank per step

```


How to read this diagram

1. **The MFU rows** multiply 6 by the parameter count and the measured tokens per second, then divide by 8 times the H100 BF16 peak. The 50.0% row is the health bar for a 70B run on 8 H100s. A measured 21% means something is badly wrong; the profiling rules in the track's Chapter 5 say where to look.
2. **The bubble rows** apply $(p-1)/(m+p-1)$. Eight stages with 32 microbatches wastes 18% of the schedule; with 8 microbatches it wastes 47%. The lesson is one inequality: keep microbatches well above stage count.
3. **The FSDP rows** show the $2(N-1)/N$ factor saturating: 64 ranks and 512 ranks move nearly the same bytes per rank (551 vs 559 GB). At 100 GB/s effective interconnect that is over 5 seconds of pure communication per step. Large runs add tensor parallelism to move some of that traffic onto NVLink.

Throughput and MFU sized the run. Now what the network costs: ring versus tree, the alpha-beta model, and the crossover point.

MFU: achieved model FLOPs over hardware peak (70B, 8x H100, BF16)

MFU counts only useful math; comm, recompute, and bubbles are not in the numerator



MFU = $(6 \times N \times \text{tok/s} / n_{\text{gpus}}) / \text{peak}$. A healthy 70B run on 8 H100s beats 9,423 tok/s (50%). A measured 21% means something is badly wrong.

MFU = $(6 \times N \times \text{tok/s} / n_{\text{gpus}}) / \text{peak}$. The 9,423 tok/s row lands exactly on 50%: the health bar a 70B run on 8 H100s should beat. A measured 21% means something is badly wrong; below 40% means profile now.

What to look at first: the green healthy band from 50 to 65%. The red bar at 21.2% sits in the profile-now zone. The one takeaway: MFU counts only useful math, so comm, recompute, and bubbles never appear in the numerator.

2.2 Collectives and the fabric

#	Quantity	Value	Units / meaning
10	Ring all-reduce bytes	$2(N-1)/N * M$	Bytes moved per rank for message size M on N ranks. About 2M at large N. Per-rank traffic barely depends on N: 8 to 64 ranks raises the factor from 1.75 to 1.97.
11	All-to-all bytes	$(N-1)/N * M$	About M per rank. An all-reduce costs roughly twice an all-to-all on the same tensor (scatter plus gather).
12	Alpha-beta model	$T = \text{steps} * \alpha + \text{bytes}/\text{BW}$	Seconds. Ring: $2(N-1)$ steps of M/N bytes. Tree: $2 * \log_2(N)$ phases of M bytes. Alpha is per-step latency, BW effective link bandwidth.
13	Ring/tree crossover	~2.3 MB	On 64 ranks, 40 GB/s IB, 5 us per step. Tree wins below, ring wins above. Recompute for your fabric with the calculator.
14	NVLink 4 (H100)	900 GB/s	Bidirectional, per GPU, inside one 8-GPU node via NVSwitch.
15	InfiniBand NDR400	50 GB/s	Per link (400 Gb/s), between nodes. The 900/50 ratio is the 18x cliff: crossing the node boundary costs 18x in per-GPU bandwidth.
16	PCIe 5.0 x16	~64 GB/s	GPU-to-host, for reference when sizing H2D traffic.

17	TP all-reduce worked	17.5 GB -> 19 ms vs 350 ms	17.5 GB of ring traffic per rank. On NVLink: 19 ms. On IB: 350 ms. This single calculation is why TP stays inside the node.
18	DP grad sync worked	7B, 64 ranks: 0.55 s	27.6 GB per rank over IB at 50 GB/s. Against a 2 s step, comm is 27%: overlap with backward is mandatory, not optional.
19	H100 memory bandwidth	~3,350 GB/s	HBM3 per H100 SXM. A100 80 GB is 2,039 GB/s. Useful when a kernel is memory-bound and you need the roofline ceiling.
20	ZeRO stages (7B example)	Z1: 38.5, Z2: 26.3, Z3: 14.0 GB/GPU	Bytes per GPU: Z1 = $2 + 2 + 12/8$, Z2 = $2 + 14/8$, Z3 = $16/8$ plus one live layer. Comm: Z1 ~1.75M, Z3 ~2.6M (about 1.5x DDP).

$$T = \text{steps} \times \alpha + \text{bytes} / \text{BW}$$

Where: T is the collective time in seconds. Steps is the algorithm's step count (ring: $2(N-1)$, tree: $2 \cdot \log_2(N)$). α (alpha) is the per-step latency. Bytes is the data moved per step. BW is the effective link bandwidth. Ring: $2(N-1)$ steps of M/N bytes. Tree: $2 \cdot \log_2(N)$ phases of M bytes.

Worked on 64 ranks with 40 GB/s IB and 5 μ s per step. Ring time = $126 \cdot 5 \mu\text{s} + 126 \cdot (M/64) / 40 \text{ GB/s}$. Tree time = $12 \cdot 5 \mu\text{s} + 12 \cdot M / 40 \text{ GB/s}$. They are equal at $M \approx 2.3 \text{ MB}$. Tree wins below, ring wins above (table row 13).

Ring all-reduce, 4 ranks, message split into 4 chunks (A B C D):

step 1 (reduce-scatter):	step 2:	step 3:	step 4 (all-gather):
rank0 --A--> rank1 > rank1	rank0 --B--> rank1	rank0 --C--> rank1	rank0 --D--> rank1
rank1 --B--> rank2 > rank2	rank1 --C--> rank2	rank1 --D--> rank2	rank1 --A--> rank2
rank2 --C--> rank3 > rank3	rank2 --D--> rank3	rank2 --A--> rank3	rank2 --B--> rank3
rank3 --D--> rank0 > rank0	rank3 --A--> rank0	rank3 --C--> rank0	rank3 --D--> rank0

(each rank adds what it receives to its own chunk)

Each rank sends $2*(N-1)$ chunks of size $M/N = 2*3/4*M$ bytes total.

That is the $2(N-1)/N$ factor: two trips around the ring, one per phase.

How to read this diagram

1. **Cut the buffer into N chunks.** Rank 0 owns chunk 0, rank 1 owns chunk 1, and so on. Each rank is responsible for summing one chunk.
2. **Steps 1-3: reduce-scatter.** Every rank sends its current chunk for the neighbor's owned index rightward and adds the received chunk to its own. After N-1 steps, each rank holds the fully summed chunk it owns.
3. **Step 4 onward: all-gather.** The summed chunks circulate N-1 more steps, copied not added. After N-1 steps, every rank holds all summed chunks: the full result.
4. **Count the bytes.** Each rank moves $2(N-1)$ chunks of size M/N . The ring keeps every link busy in every step, which is why it wins on large messages.

```

import math

# Alpha-beta model: ring vs tree all-reduce across message sizes.
# WHY: this is the trade-off NCCL's tuner navigates on every collective.
#   When NCCL picks the loser, this model proves it.
# WHAT BREAKS: alpha and bandwidth are EFFECTIVE values, not spec sheets.
#   A sick link raises alpha ~10x and the model will disagree with the
#   tuner. That disagreement is itself a diagnostic signal.

def ring_allreduce_time(m_bytes, n_ranks, alpha_s, bw_bytes_s):
    # Ring: 2(N-1) steps, each moving M/N bytes. Bandwidth-optimal.
    steps = 2 * (n_ranks - 1)
    chunk = m_bytes / n_ranks
    return steps * alpha_s + steps * chunk / bw_bytes_s

def tree_allreduce_time(m_bytes, n_ranks, alpha_s, bw_bytes_s):
    # Tree: 2*log2(N) phases, each moving the full M bytes.
    # Latency-optimal: step count grows with log N, not N.
    phases = 2 * math.log2(n_ranks)
    return phases * alpha_s + phases * m_bytes / bw_bytes_s

def crossover_point(n_ranks, alpha_s, bw_bytes_s):
    # Binary search for the message size where the times are equal.
    lo, hi = 1.0, 1e12
    for _ in range(80):
        mid = (lo + hi) / 2
        if tree_allreduce_time(mid, n_ranks, alpha_s, bw_bytes_s) < \
            ring_allreduce_time(mid, n_ranks, alpha_s, bw_bytes_s):
            lo = mid
        else:
            hi = mid
    return (lo + hi) / 2

# Fabric: inter-node InfiniBand, effective 40 GB/s, 5 us per step.
N, ALPHA, BW = 64, 5e-6, 40e9
print(f"{'message':>10} {'ring':>12} {'tree':>12}  winner")
for m in (4e3, 1e6, 64e6, 1e9):
    tr = ring_allreduce_time(m, N, ALPHA, BW)
    tt = tree_allreduce_time(m, N, ALPHA, BW)
    win = "tree" if tt < tr else "ring"
    print(f"{m:>9.0f}B {tr*1e6:>10.1f}us {tt*1e6:>10.1f}us  {win}")
cross = crossover_point(N, ALPHA, BW)
print(f"crossover: tree wins below ~{cross/1e6:.1f} MB, ring wins above")

```

message	ring	tree	winner
4000B	630.2us	61.2us	tree
1000000B	679.2us	360.0us	tree
64000000B	3780.0us	19260.0us	ring
1000000000B	49848.8us	300060.0us	ring

crossover: tree wins below ~2.3 MB, ring wins above


How to read this diagram

1. At **4 KB** the tree finishes in 61 us while the ring takes 630 us: latency dominates and the log-step tree wins by 10x.
2. At **1 GB** the ring takes about 50 ms while the tree takes 300 ms: bandwidth dominates and the ring wins by 6x.
3. **The crossover near 2.3 MB** is where the two curves meet on this fabric. Map it to training. DDP gradient buckets are typically tens of MB (ring territory). TP per-layer all-reduces often run KB to a few MB (tree territory). That is why NCCL picks the tree there.
4. **If you force Ring everywhere**, the small collectives get slower. The tuner is usually right. Override it only with a measurement in hand.

You pay for the network. Now failure math: the Daly interval, how MTBF scales with GPU count, and what a Llama-3-scale run really pays.

Ring vs tree all-reduce: latency wins small, bandwidth wins large

64 ranks, 40 GB/s IB, 5 us per step; ring = $2(N-1)$ steps, tree = $2 \cdot \log_2(N)$ phases



Map to training: DDP gradient buckets are tens of MB (ring territory); TP per-layer all-reduces are KB to a few MB (tree territory). Recompute the crossover for your fabric.

On 64 ranks with 40 GB/s IB and 5 us per step, the tree finishes a 4 KB all-reduce in 61 us. The ring takes 630 us. At 1 GB the ring takes about 50 ms while the tree takes 300 ms. DDP gradient buckets live in ring territory; TP per-layer all-reduces live in tree territory.

What to look at first: the green dashed crossover line near 2.3 MB. Left of it the amber tree line is lower; right of it the blue ring line is lower. The x axis is logarithmic: each tick multiplies the message size. The one takeaway: recompute this crossover for your own fabric instead of trusting defaults.




Figure 4.8. The protocol rule in one picture. Below about 2.3 MB, latency dominates: the tree with LL or LL128 wins. Above it, bandwidth dominates: the ring with Simple wins. DDP gradient buckets live on the right; TP per-layer all-reduces live on the left.

How to read this diagram

1. Read the x-axis as message size on a log scale: each tick multiplies the bytes.
2. The dashed line at about 2.3 MB is the crossover on this fabric (64 ranks, 40 GB/s, 5 us per step).
3. Left of the line, the tree wins because its 12 phases beat the ring 126 steps when latency dominates. Right of the line, the ring wins because it saturates bandwidth.
4. What it is: the alpha-beta model as a map. Why it exists: NCCL autotunes to this map, and overrides need a measurement. What breaks: trusting defaults on a different fabric (recompute the crossover), or forcing LL on large messages (lower bandwidth).




Figure 4.9. Why bucket size matters. Small buckets let each all-reduce overlap with the backward compute of later layers. One giant bucket cannot overlap anything: the whole all-reduce sits exposed. Exposed communication is the fraction not hidden behind compute.

How to read this diagram

1. Read the top green bar as the backward pass running left to right.
2. The blue blocks are per-bucket all-reduces. Each one runs while the backward pass continues: hidden time.
3. The red block is one giant bucket: it starts only after the last gradient is ready, so nothing hides it.
4. What it is: the overlap argument for small buckets. Why it exists: communication you hide costs nothing. What breaks: buckets so small that latency dominates (the 2.3 MB rule again), or a trace with no overlap (fix bucket size or microbatch size first).

2.3 Checkpointing and failure math

#	Quantity	Value	Units / meaning
21	Checkpoint anatomy (70B)	700 GB total	140 GB parameters (70B x 2 bytes BF16) + 560 GB Adam state (70B x 8 bytes: two FP32 states). Optimizer dominates 4:1. RNG + dataloader state are MBs but required for exact resume.
22	Daly interval	$\tau^* = \sqrt{2 \cdot C \cdot \mu}$	Seconds. C = save stall, μ = MTBF. Example: C = 200 s, μ = 50.7 h -> 142 min. With async (C = 30 s) -> 55 min.
23	MTBF scaling	Inverse with GPU count	Per-GPU-hour failure rate $\sim 1.97 \times 10^{-5}$ (from Llama 3: 419 unexpected / (16,384 x 54 x 24)). 1,000 GPUs -> ~ 50.7 h. 16,384 GPUs -> ~ 3 h.
24	Llama 3 run figures	16,384 H100, 54 days, 466 interruptions	47 planned, 419 unexpected: about one failure every 3 hours. GPU problems 58.7% of unexpected. Effective training time stayed above 90% via automation.
25	Recovery budget	< 10 min total	Detect + isolate + resume. From the guardian-loop lab: fast recovery (2 min/failure) held 86.6% effective time vs 61.0% for slow recovery (10 min/failure) at the same failure density.

26	Autoretry payoff	33.3% chain success	2.7x better than manual recovery. Median 11-min retry interval; downtime 1.9 h vs 3.3 h manual. From a public study of a ~500-GPU cluster.
27	Fail-slow cost	> 1 h silent, minutes to fix	Average fail-slow episode runs over an hour before anyone notices; remediation is 30 s (rebalance) to a few minutes (checkpoint and restart on a healthy node).
28	Spike detector	median + 10*MAD, 5 steps	Loss above median + 10x median-absolute-deviation for 5 consecutive steps = sustained spike. One step is noise; five is a signal.

```

import math

# Checkpoint budget: byte anatomy, Daly interval, restore times.
# WHY: these three numbers set the run's failure policy. Save too rarely
#   and each failure wastes hours; save too often and the write tax eats
#   the run. Daly's formula is the optimum, not a guess.
# WHAT BREAKS if changed: effective filesystem bandwidth is not the spec
#   sheet. Measure per-rank bandwidth on YOUR cluster during bring-up.

def checkpoint_bytes(n_params, bytes_per_param=2, optimizer_factor=4.0):
    # Params at 2 bytes (BF16). Adam keeps two FP32 states per parameter
    # (momentum + variance): 8 bytes = 4x the parameter bytes.
    params = n_params * bytes_per_param
    return params * (1.0 + optimizer_factor)

def daly_interval_s(mtbfs, checkpoint_cost_s):
    # sqrt(2*C*mu). Balances expected lost work (grows with the interval)
    # against write overhead (shrinks with the interval). Assumes
    # failures are memoryless (Poisson-ish), true for hardware faults.
    return math.sqrt(2.0 * checkpoint_cost_s * mtbfs)

def restore_time_s(total_bytes, n_ranks, per_rank_read_bw,
                   fixed_overhead_s=60):
    # Sharded restore: every rank reads only its own shard in parallel.
    return total_bytes / n_ranks / per_rank_read_bw + fixed_overhead_s

TOTAL = checkpoint_bytes(70e9)
print(f"checkpoint size (70B, Adam, BF16): {TOTAL/1e9:.0f} GB")

# Save cost: 8 ranks sustaining 1 GB/s each to the parallel filesystem.
save_cost = TOTAL / (8 * 1e9)
print(f"sharded write time: {save_cost:.0f} s ({save_cost/60:.1f} min)")

# Daly with the 1k-GPU MTBF of 50 hours from the Llama 3 failure scaling.
interval = daly_interval_s(50 * 3600, save_cost)
print(f"Daly interval (C={save_cost:.0f}s, MTBF=50h): {interval/60:.0f} min")

# Async checkpointing cuts the stall to the staging copy (~30 s).
interval_async = daly_interval_s(50 * 3600, 30.0)
print(f"Daly interval (C=30s async, MTBF=50h): {interval_async/60:.0f} min")

# Restore: 8 ranks at 5 GB/s each; compare with one rank reading all.
print(f"sharded restore (8 ranks): {restore_time_s(TOTAL, 8, 5e9)/60:.1f} min")
print(f"single-writer restore: {(TOTAL/5e9 + 60)/60:.1f} min (do not do this)")

```

```

checkpoint size (70B, Adam, BF16): 700 GB
sharded write time: 88 s (1.5 min)
Daly interval (C=88s, MTBF=50h): 94 min
Daly interval (C=30s async, MTBF=50h): 55 min
sharded restore (8 ranks): 1.3 min
single-writer restore: 3.3 min (do not do this)

```

How to read this diagram

1. **700 GB** is the checkpoint, and the optimizer state is the reason. Any scheme that saves only parameters is not a resume; it is a restart with amnesia.
2. **94 minutes** is the Daly optimum for an 88-second sharded write at 50-hour MTBF. Async checkpointing cuts the stall to the staging copy and the optimum drops to 55 min.
3. **The restore rows** show why sharding matters on both directions: 8 ranks reading in parallel finish in 1.3 min. The single-writer number looks modest at 8 ranks, but its real cost appears at larger rank counts and in the redistribution step this model omits.






Figure 4.11. Restore is a read problem. Sharded restore lets each of the 8 ranks read its own shard: about a minute for 700 GB. A single-reader design serializes the whole checkpoint through one rank, then pays redistribution on top: about 40 minutes. Shard the writes and the reads.

The GPU numbers are in place. What changes on TPU and XLA, and what transfers.

Failure math: MTBF falls as GPU count rises; the checkpoint is mostly optimizer state

per-GPU-hour failure rate $\sim 1.97\text{e-}5$ from the Llama 3 run (419 unexpected / 16,384 x 54 x 24)

MTBF vs fleet size (inverse scaling)**Checkpoint anatomy, 70B + Adam + BF16 (700 GB total)**

Any scheme that saves only parameters is not a resume; it is a restart with amnesia.

Llama 3 scale: 16,384 H100s, 54 days, 466 interruptions (47 planned, 419 unexpected): about one failure every 3 hours. Effective training time stayed above 90% via automation.

Per-GPU-hour failure rate of about $1.97\text{e-}5$ (from the Llama 3 run) means 1,000 GPUs see ~ 50.7 h MTBF while 16,384 GPUs see about 3 h. The checkpoint is 700 GB because Adam state dominates 4:1 over parameters; RNG and dataloader state are megabytes but required for exact resume.

What to look at first: the two blue MTBF bars and how fast the second shrinks. Then the stacked bar: the purple optimizer segment is four times the blue parameter segment. The one takeaway: any scheme that saves only parameters is a restart with amnesia, not a resume.

2.4 The TPU and XLA numbers worth keeping

#	Quantity	Value	Units / meaning
29	TPU v6e (Trillium)	~918 BF16 TFLOPS, 32 GB HBM3	Per chip. 256-chip 2D torus pods (a torus is a mesh network where each chip connects to its neighbors and the edges wrap around, so every chip has equal standing and there are no fast islands): $256 \times 32 \text{ GB} = 8 \text{ TB}$ of HBM in one pod.
30	TPU v7 (Ironwood)	4,614 FP8 TFLOPS, 192 GB HBM3e	Per chip. 9,216-chip 3D torus pod: 42.5 FP8 exaflops. First TPU with native FP8.
31	Pod-scale sync cost	671B FP8: ~6.7 s per full all-reduce	2 x 671 GB of gradients at ~200 GB/s per chip of ICI. The lesson: at pod scale you never sync full parameters every step; you shard them.
32	XLA compile key	one binary per (shape, dtype)	Change the sequence length and you pay a recompile. Bucketed padding (128, 256, 512) is the standard payment.

What transfers from the GPU world. Mesh = named process groups. PartitionSpec = sharding spec. The same collectives (all-reduce, all-gather, all-to-all) run on ICI (the TPU's Inter-Core Interconnect) instead of NVLink/InfiniBand. What does not transfer: the TP-in-side-the-node rule (a torus has no NVLink islands; the cost model is hop count), bfloat16-first numerics with hardware-specific overflow behavior, and host-driven execution with a different profiler.

Key takeaways

- 32 numbers, each with units and an operational meaning. The ones to say from memory: 6ND, MFU 50-65%, bubble $(p-1)/(m+p-1)$, $2(N-1)/N$, Daly $\sqrt{2C_{mu}}$, 700 GB for 70B + Adam. Also: MTBF inverse with GPU count, and the $\sim 18\times$ NVLink-vs-InfiniBand per-layer collective gap from 1.1.
- The text marks site-specific values (filesystem bandwidth, IB latency, cloud failure rates). Measure them during bring-up and substitute.

Next: Section 3, six traps with symptoms first, so you can match what you see.

TPU and XLA numbers worth keeping

same collectives, different cost model: the torus prices hops, not NVLink islands

TPU v6e (Trillium)

~ 918 BF16 TFLOPS, 32 GB HBM3 per chip; 256-chip 2D torus pod = 8 TB HBM in one pod

TPU v7 (Ironwood)

4,614 FP8 TFLOPS, 192 GB HBM3e per chip; 9,216-chip 3D torus pod = 42.5 FP8 exaflops; first TPU with native FP8

pod-scale sync lesson

671B FP8: ~ 6.7 s per full all-reduce at ~ 200 GB/s ICL. At pod scale you never sync full parameters every step; you shard them.

What transfers: mesh = named process groups, PartitionSpec = sharding spec. What does not: the TP-inside-the-node rule, bfloat16-first numerics, host-driven execution and its profiler.

TPU v6e: ~ 918 BF16 TFLOPS and 32 GB HBM3 per chip, 256-chip 2D torus pods. TPU v7 (Ironwood): 4,614 FP8 TFLOPS and 192 GB HBM3e per chip, the first TPU with native FP8. At pod scale a full 671B FP8 all-reduce costs ~ 6.7 s, so you never sync full parameters every step: you shard them.

What to look at first: the two cards side by side, then the amber lesson box. The one takeaway: the torus prices hop count, not NVLink islands, so the TP-inside-the-node rule does not transfer.

Section 3: Classic traps

The math is right and the run is still red. Section 3 is the symptom-first field guide: stragglers, NCCL hangs, checkpoints that are not resumes, misleading FLOPS, metadata storms, and profiling the wrong thing.

Each trap names its symptom first, so you can match what you see. The fix follows.

Trap 1: the straggler that taxes every step

Symptom. The run is slower than the arithmetic says it should be. Loss looks normal. Average GPU utilization looks fine. Nothing is on fire; everything is just slow.

What is wrong. Synchronous training finishes each step when the slowest rank finishes. One rank at 2.6 s while 999 ranks run at 2.0 s makes every step 2.6 s. A 30% tax from one GPU in a thousand. The average hides it because 999 healthy ranks dilute one sick one.

The four kinds. Compute straggler (thermal throttling, degraded card), input straggler (slow storage path on one rank), H2D straggler (host-to-device transfer slowed by NUMA misplacement or missing pinned memory; NUMA is multi-socket servers where memory attached to the other CPU socket is slower, so a GPU process bound to the wrong socket pays a transfer tax), network straggler or fail-slow (a degraded NIC that delivers every byte, just slowly). Fail-slow is the nastiest. The NCCL timeout never fires because the job still makes progress. Community reports put the average episode at over an hour of silent slowdown.

Detection. Compare peers, not thresholds. Log per-rank step times, take the median per rank over a window, and flag ranks slower than the fleet median by the MAD scale. Straggler factor = max/median: above 1.1x investigate, above 1.3x act. Separate compute slowness from collective slowness with NCCL per-collective timing: slow inside collectives points at the NIC, slow in kernels points at the GPU.

```

import random
import statistics

# Straggler detection on synthetic per-rank step times.
# WHY: the straggler factor is the cheapest high-value diagnostic in the
#   role. One slow node can cost half the cluster's throughput, and the
#   fix (quarantine the node) takes minutes once you know which one.
# WHAT BREAKS: the 1.1x threshold assumes steady-state steps. Warm-up,
#   checkpoint, and eval steps are legitimately slower and must be
#   excluded first, or every rank looks slow once per hour. Never put a
#   barrier before the timestamps: barriers hide the skew you measure.

def find_stragglers(step_times, threshold=1.10):
    # median, not mean: one terrible rank must not drag the baseline.
    med = statistics.median(step_times)
    factor = max(step_times) / med
    flagged = [(i, t) for i, t in enumerate(step_times)
               if t > threshold * med]
    return factor, flagged, med

# 64 ranks at ~2.0 s. Rank 17 is thermally throttling; rank 52 has a
# fail-slow NIC (slower AND more variable).
rng = random.Random(11)
times = [rng.gauss(2.0, 0.05) for _ in range(64)]
times[17] = rng.gauss(2.35, 0.05)
times[52] = rng.gauss(3.05, 0.08)

factor, flagged, med = find_stragglers(times)
print(f"median step {med:.3f}s, straggler factor {factor:.2f}x")
for rank, t in flagged:
    print(f"  rank {rank}: {t:.2f}s ({t/med:.2f}x median)")
drag = max(times) - med
print(f"drag per step: {drag:.2f}s, about {drag/med:.0%} of step time wasted")

```

```

median step 2.006s, straggler factor 1.54x
  rank 17: 2.38s (1.19x median)
  rank 52: 3.09s (1.54x median)
drag per step: 1.08s, about 54% of step time wasted


```

How to read this diagram

1. **Rank 52 is the fire.** At 1.54x the median it forces all 63 other ranks to wait an extra 1.08 s every step. Over half the step time wasted cluster-wide.
2. **Rank 17 is worth watching** at 1.19x (likely thermal throttling) but is not the fire today.
3. **The median-based rule matters.** Outliers would pull a mean up and hide them. Run this on a rolling window of step times, exclude warm-up and checkpoint steps, and feed flagged ranks to the quarantine list.

One slow rank taxes every step

per-rank step times, 8 ranks, median 2.0 s



Synchronous training finishes each step when the slowest rank finishes: r4 at 3.1 s (1.54x) makes all 8 ranks wait.

One slow rank taxes every step. Training is synchronous: each step finishes when the slowest rank finishes, so rank 4 at 3.1 s (1.54x the 2.0 s median) makes all 8 ranks wait. The straggler factor is max step time over median; above 1.1x investigate, above 1.3x act.

Response, in order of cost. Rebalance microbatches (give the slow rank proportionally less work; a 30-second fix when the imbalance is stable). Evacuate and restart on a spare (minutes with fast multi-tier checkpoints). Mark the node unhealthy so the scheduler stops handing it out.

Trap 2: the NCCL hang that was never the network

Symptom. The job stops making progress. GPUs sit at 0% utilization. Nothing prints an error. The logs just stop.

What is happening, in likelihood order:

1. A collective inside a conditional branch that not all ranks take: rank 0 logs a metric with an all-reduce inside `if rank == 0` and waits forever.
2. Uneven data: one rank's loader runs out of batches, skips its last backward and its all-reduce, while the others wait.
3. One rank genuinely slow or sick.
4. The network. The network is last on a healthy cluster, not first.

The debugging order.

```

job stops progressing
are GPUs idle AND logs stopped?
  no -> not a comm hang (check dataloader, checkpoint write)
  yes ->
set NCCL_DEBUG=INFO, rerun
do all ranks log the same collective sequence?
  no -> the rank that diverged is the bug:
        different op/size/dtype = code bug
        rank never reached the op = stall upstream (dataloader)
  yes ->
run nccl-tests all_reduce_perf on the same node set
does the isolated test hang too?
  no -> bug is in your code's collective pattern
  yes -> bisect the node set; hang follows nodes ->
        hardware or fabric (IB link flaps, GPU Xid errors: Xid codes are
        NVIDIA's numbered GPU error reports)

```

How to read this flowchart. Each level rules out a cheaper cause before you pay for a more expensive one. Code divergence first (free to check), dataloader second, slow hardware third, fabric last. Reproduce with 2 ranks and a tiny tensor before debugging a thousand-GPU hang. If the repro hangs, the bug is in your launch or code, and you iterate in seconds instead of hours.


The environment knobs, in the order to try them:

1. **NCCL_DEBUG=INFO** (then **NCCL_DEBUG_SUBSYS=ALL** for the full firehose).
2. **NCCL_TIMEOUT**: shorten while debugging so hangs fail fast; the default can be 30 minutes, which turns every hang into a 30-minute mystery.
3. **NCCL_ALGO** / **NCCL_PROTO**: force Tree/LL for small-message hangs, Ring/Simple for large.
4. **NCCL_MIN_NCHANNELS** / **NCCL_MAX_NCHANNELS** and **NCCL_SOCKET_PERTHREAD** when a sick path needs sidestepping.
5. **The flight recorder** (**TORCH_NCCL_TRACE_BUFFER_SIZE**, **TORCH_NCCL_DUMP_ON_TIMEOUT=1**, **TORCH_NCCL_DESYNC_DEBUG=1**) names the guilty rank directly.

The mismatched-env-var variant. NCCL reads its configuration from the environment on every rank. If ranks disagree (different `NCCL_ALGO` on different nodes because a launch script set it on half the job, or `NCCL_SOCKET_IFNAME` pointing at different interfaces), collectives can hang or silently take different algorithms. The fix is unglamorous: dump the NCCL-relevant environment on every rank at startup and diff them. A config that only diverges on the nodes you added last often causes the hang that only appears past 2,000 GPUs.

The NCCL hang that was never the network: debug in likelihood order

each level rules out a cheaper cause before you pay for a more expensive one



When the job stops with idle GPUs and silent logs, start here. Check collective divergence with `NCCL_DEBUG=INFO` first (free), then the dataloader, then isolate with `nccl-tests`, and bisect the node set last. The network is last on a healthy cluster, not first.

What to look at first: the top diamond, then each no-branch on the left in order. The arrows only move down on yes. The one takeaway: reproduce with 2 ranks and a tiny tensor before debugging a thousand-GPU hang. If the repro hangs, the bug is in your launch or code.

Trap 3: the checkpoint that was not a resume

Symptom. After a preemption or crash, the run restarts from the checkpoint. The loss jumps, or training silently continues on a wrong data order, or the restore fails outright.

Four variants, each a real bug.

1. **The missing optimizer state.** The checkpoint saved the 140 GB of parameters but not the 560 GB of Adam state. The resume is not a resume; it is a restart with amnesia, and the loss curve pays for it.
2. **The single retained checkpoint.** One checkpoint, one write, one corruption: everything is lost. Keep the last 3 plus periodic older ones. A checkpoint you have never restored is a rumor: verify by actually restoring into a fresh process and comparing a parameter hash.
3. **The async staging race.** Two background writes overlap because the code did not block on the in-flight future, and the second staging copy clobbers the first. The rule is absolute: exactly one in-flight write; `future.result()` before the next save. Write errors surface at `.result()`, not at save time: check it or failures stay silent.

4. **The preemption mid-write.** A node dies between the local NVMe write (NVMe: the fast local SSD on each node) and the copy-out to durable storage. The run guardian must track which checkpoints are durable and which are still local, and the copy-out must be prompt.

The resume checklist:

- Weights hash matches.
- Optimizer state loaded, not re-initialized.
- RNG states restored.
- Data loader resumes at the right sample.
- LR schedule resumes at the right step. Schedules index by step: resuming weights at step 5,000 with a scheduler at step 0 silently retrains with the wrong schedule.
- The loss curve continues smoothly with no jump.

What checkpointing cannot do. It preserves state faithfully, including corrupted state: a GPU that computes wrong results without raising an error gets its corruption faithfully saved. It does not fix poisoned data or a bad recipe. Detection picks which checkpoint is worth restoring; the checkpoint is only as good as that judgment.

```
time ----->

GPU thread:  | T0 | T1 | stage | T3 | T4 | T5 | T6 |   training steps
              (copy state to pinned CPU memory;
              ~200 ms: the only pause)

bg thread:    |<----- write to storage ----->|
              (overlaps T3-T6)

future:       [ exactly one in-flight write:
                the next checkpoint blocks here ]
```

How to read this diagram

1. **Steps T0-T1 run normally.** Nothing about checkpointing is visible.
2. **The staging pause** copies the state dict to pinned CPU memory. Training stops for a fraction of a second. Pinned memory matters because the storage layer DMAs (direct memory access: hardware moving bytes without the CPU) from it without the OS paging it out mid-write.
3. **The background thread writes** the staged copy to the filesystem while steps T3-T5 run at full speed. The background thread must not touch GPU memory.
4. **The next checkpoint waits** on the in-flight future. One in-flight write only, or the staging buffer gets clobbered. If the filesystem stalls, the fix is backpressure with a clear signal, not a bigger buffer.

The checkpoint that was not a resume: four variants, one checklist

a checkpoint you have never restored is a rumor: verify by restoring into a fresh process

Four variants, each a real bug:

- missing optimizer state**
saved 140 GB of params, not the 560 GB of Adam state
- single retained checkpoint**
one write, one corruption: everything lost; keep K=3 plus daily
- async staging race**
two background writes overlap; exactly one in-flight write
- preemption mid-write**
node dies between local NVMe write and copy-out to durable storage

The resume checklist:

- | | |
|---|--|
| ☒ weights hash matches | ☒ optimizer state loaded, not re-initialized |
| ☒ RNG states restored | ☒ dataloader resumes at the right sample |
| ☒ LR schedule resumes at the right step | ☒ loss curve continues smoothly, no jump |

What checkpointing cannot do: it preserves state faithfully, including corrupted state. Detection picks which checkpoint is worth restoring; the checkpoint is only as good as that judgment.

The four real bugs: missing optimizer state, a single retained checkpoint, the async staging race, and preemption mid-write. The checklist: weights hash matches, optimizer state loaded, RNG restored, dataloader at the right sample, LR schedule at the right step, loss continues smoothly.

What to look at first: the four colored variant rows, then the green checklist. The one takeaway: checkpointing preserves state faithfully, including corrupted state, so detection picks which checkpoint is worth restoring.

Trap 4: peak FLOPS is not achieved FLOPS

Symptom. Two runs report MFUs that cannot both be true. Your MFU disagrees with a published number by a factor of two, or a config change "improves MFU" while tokens per second stays flat.

The three ways the ratio lies.

- Wrong peak for the precision.** The H100 BF16 dense peak is 989 TFLOPS; the FP8 peak is roughly double. Quoting BF16 math against the FP8 peak halves the MFU silently. Always state the peak: "55% against the 989 TFLOPS BF16 dense peak."
- Wrong units.** TFLOPS vs FLOPS is a factor of $1e12$. It happens in spreadsheet math more often than anyone admits.
- Counting the wrong FLOPs.** MFU counts only useful model math. Communication, recomputation from gradient checkpointing (it drops activations during the forward pass and recomputes them in backward: less memory, more compute), and idle bubbles are not in the numerator. A run with heavy recomputation can show lower MFU while being the right engineering trade (it bought a bigger batch or a longer context).

The honest reading. MFU measures efficiency; goodput measures delivered value. A 60% MFU run with constant restarts can deliver fewer useful tokens than a 50% MFU run that never stops. Quote both, and quote the peak.

Peak FLOPS is not achieved FLOPS: three ways the ratio lies

MFU measures efficiency; goodput measures delivered value: quote both, and quote the peak

Three ways the ratio lies:

wrong peak for the precision

BF16 math quoted against the FP8 peak halves the MFU silently

wrong units

TFLOPS vs FLOPS is a factor of $1e12$; it happens in spreadsheet math

counting the wrong FLOPs

comm, recompute, and idle bubbles are not useful model math

honest reading

a 60% MFU run with constant restarts can deliver fewer useful tokens than a 50% MFU run that never stops

The honest line: "55% against the 989 TFLOPS BF16 dense peak." Two runs reporting MFUs that cannot both be true means at least one peak, unit, or numerator is wrong.

The ratio lies three ways: the wrong peak for the precision (BF16 math against the FP8 peak halves MFU silently), wrong units (TFLOPS vs FLOPS is $1e12$), and counting comm, recompute, or bubbles as useful math. The honest line states the peak: "55% against the 989 TFLOPS BF16 dense peak."

What to look at first: the three colored lie boxes, then the green honest-reading box. The one takeaway: a 60% MFU run with constant restarts can deliver fewer useful tokens than a 50% MFU run that never stops. Quote goodput next to MFU.

Trap 5: the small-file metadata storm

Symptom. Checkpoint writes take far longer than the byte math says they should.

Bandwidth monitors show the link mostly idle while the write crawls.


What is wrong. Parallel filesystems are built for a few large files, not millions of small ones. A checkpoint format that writes one file per tensor (tens of thousands of files) spends more time in metadata operations than in data transfer. Sharded formats that write one or a few large files per rank are not just tidier; they are measurably faster.

The fix. Keep shard files large and few. This is also why the two-tier pattern (local NVMe staging, parallel filesystem durable) wins: the fast tier absorbs the write, the background copy handles durability.

The small-file metadata storm: bytes said minutes, the filesystem said hours

parallel filesystems are built for a few large files, not millions of small ones

Where the write time actually goes (one file per tensor):



the fix: few large files

sharded formats write one or a few large files per rank

two-tier pattern

local NVMe absorbs the write; background copy handles durability

One file per tensor means tens of thousands of files. Every create, open, and sync is a metadata round trip while the link sits idle. The fix: shard formats that write one or a few large files per rank. Use the two-tier pattern: local NVMe absorbs the write and a background copy handles durability.

What to look at first: the thin blue transfer bar against the long red metadata bar in the measured row. The one takeaway: keep shard files large and few.

Trap 6: profiling the wrong thing

Three misreads that waste weeks.

1. **First-step noise.** The first steps include compilation, autotuning, and allocator warm-up. Never profile step 1. Warm up for tens of steps first.
2. **The average lie.** A 50% MFU average can hide 90% on half the ranks and 10% on the other half. Always look at the distribution (the straggler analysis from Trap 1) before the mean.
3. **Profiler overhead.** Nsys itself slows the run. Compare profiles to profiles, not to unprofiled runs, when judging relative changes. Measure absolute step time without the profiler attached.

The one tuning discipline:

1. **Change one variable at a time,** in the order that usually pays: microbatch size, gradient bucket size, TP degree within the node.
2. **Touch NCCL knobs only after the layout is right.**
3. **Touch dataloader workers only if the trace shows host starvation.**
4. **Each change gets a before/after measurement** of step time and MFU, medians over enough steps, recorded with the config.
5. **Do not ship a "win" that sits inside the noise band.**

```

import numpy as np

# Loss-spike detector on a synthetic training run.
# WHY: MAD (median absolute deviation) sets the threshold instead of the
# standard deviation, because one giant spike does not inflate it.
# The persist counter separates a poisoned optimizer (roll back) from
# a noisy batch (keep training).
# WHAT BREAKS: k and persist must be tuned on a HEALTHY run's noise
# level. Raise the data noise and false positives appear; that is the
# detector saying the baseline changed, not a bug in the code.

rng = np.random.default_rng(0)
steps = np.arange(1200)
# Slowly decaying loss plus noise, then a 5-step spike at step 800.
loss = 8.0 / (1 + steps / 200.0) + rng.normal(0, 0.05, 1200)
loss[800:805] += 6.0

class SpikeDetector:
    def __init__(self, window=200, k=10.0, persist=5):
        # window: recent steps forming the "normal" baseline.
        # k: threshold in MAD units. persist: consecutive bad steps
        # before firing (one step can be noise).
        self.window, self.k, self.persist = window, k, persist
        self.hist, self.streak = [], 0

    def step(self, v):
        self.hist.append(float(v))
        if len(self.hist) > self.window:
            # Oldest drops out: the baseline must track recent training,
            # not week-old loss values at a different scale.
            self.hist.pop(0)
        if len(self.hist) < self.window:
            # Not enough data yet. Never alert before the baseline exists.
            return False
        med = float(np.median(self.hist))
        mad = float(np.median(np.abs(np.asarray(self.hist) - med)))
        if v > med + self.k * max(mad, 1e-8):
            self.streak += 1
        else:
            self.streak = 0
        return self.streak >= self.persist

    def reset(self):
        # Call after a rollback: the old baseline belongs to the spiked
        # run and would misfire on the restored weights.
        self.hist.clear()
        self.streak = 0

det = SpikeDetector()
fired_at = None
for i, v in enumerate(loss):
    if det.step(v) and fired_at is None:
        fired_at = i

```

```
print(f"spike injected at step 800 (5 steps), detector fired at step
{fired_at}")

# Blip check: a 2-step blip must NOT fire (persist=5 filters it).
loss2 = 8.0 / (1 + steps / 200.0) + rng.normal(0, 0.05, 1200)
loss2[800:802] += 6.0
det2 = SpikeDetector()
fired2 = any(det2.step(v) for v in loss2)
print(f"2-step blip fired detector: {fired2} (expect False)")
```

```
spike injected at step 800 (5 steps), detector fired at step 804
2-step blip fired detector: False (expect False)
```

How to read this diagram

1. **The baseline is a rolling window**, not the whole run. Loss drifts down during training; a fixed threshold from week one would fire constantly by week three.
2. **MAD sets the scale**. On noisy data the MAD is large and the detector is lenient. On clean data it is tight. The threshold adapts to the run.
3. **Persist separates spikes from blips**. A single bad batch recovers on its own. Five bad steps in a row means the spike may have poisoned the optimizer's momentum buffers: roll back.
4. **Reset after rollback**. The detector's memory belongs to the spiked run. Restoring the model without resetting the detector leaves a baseline trained on the failure.

Key takeaways

- Six traps: the silent straggler, the hang that was code not network, the checkpoint that was not a resume. Peak vs achieved FLOPS, the metadata storm, profiling the wrong thing.
- Each has a number attached: 1.1x/1.3x, 30-minute default timeout, $K=3$ retention, the stated peak, few large files, warm-up before profiling.
- When in doubt, suspect your code first, the hardware second, the network last.

Next: Section 4, thirty questions with worked answers spanning parallelism, collectives, checkpointing, failures, and TPU/XLA.


Profiling the wrong thing: three misreads and the one tuning discipline

change one variable at a time; do not ship a win that sits inside the noise band

Three misreads that waste weeks:

- first-step noise**
steps 1-N include compilation, autotuning, allocator warm-up; never profile step 1
- the average lie**
50% MFU average can hide 90% on half the ranks and 10% on the rest; look at the distribution
- profiler overhead**
nnsys itself slows the run; compare profiles to profiles, measure absolute step time unprofiled

The tuning order that usually pays:



Each change gets a before/after measurement of step time and MFU, medians over enough steps, recorded with the config.

The misreads: first-step noise (never profile step 1), the average lie (look at the distribution, not the mean), and profiler overhead (compare profiles to profiles). The order: microbatch size, gradient bucket size, TP degree within the node, then NCCL knobs, then dataloader workers only if the trace shows host starvation.

What to look at first: the three colored misread rows, then the numbered left-to-right tuning order. The one takeaway: each change gets a before/after measurement of step time and MFU, medians over enough steps, recorded with the config. Do not ship a win inside the noise band.

Section 4: Rapid-fire Q&A

Work the arithmetic before you reveal. These are the questions interviewers ask and dashboards answer.

Thirty questions with worked answers. Cover the answer, work the arithmetic, then check. The questions span parallelism, collectives, checkpointing, failure taxonomy, profiling, and the TPU/XLA stack.

Q1. A 13B model trains on 32 H100 GPUs at 38,000 tokens/sec (global). What is the MFU? Use 989 TFLOPS peak per GPU.

- A. 5%
- B. 9%
- C. 30%
- D. 60%

Answer: B. Achieved = $6 \times 13e9 \times 38,000 = 2.96e15$ FLOPs/s. Peak = $32 \times 989e12 = 3.16e16$. MFU = $2.96e15 / 3.16e16 = 9.4\%$, about 9%. A healthy run should beat 40%; this one needs profiling immediately.

- **A, wrong:** too low; check the arithmetic before trusting a gut number.
- **C, wrong:** roughly what 120,000 tok/s would give. The throughput is not there.
- **D, wrong:** would need about 250,000 tok/s on this hardware.

Q2. Pipeline parallelism with 6 stages and 18 microbatches. What fraction of the schedule is bubble?

- A. 10%
- B. 22%
- C. 33%
- D. 45%

Answer: B. Bubble = $(p-1)/(m+p-1) = 5/23 = 21.7\%$, about 22%.

- **A, wrong:** too optimistic; forgets the drain phase also idles stages.
- **C, wrong:** roughly the answer for 6 stages and 10 microbatches.
- **D, wrong:** far too high for 18 microbatches.

Q3. A 64-rank all-reduce moves a 1 GB message on 40 GB/s links with 5 us per-step latency. Which algorithm wins, and roughly by how much?

- A. Tree, by about 2x
- B. Ring, by about 6x
- C. Ring, by about 50x
- D. They tie; message size does not matter

Answer: B. Ring ~50 ms, tree ~300 ms. The ring wins by about 6x because bandwidth dominates at 1 GB.

- **A, wrong:** the tree wins only for small messages (below ~2.3 MB here).
- **C, wrong:** 50x overstates it; that ratio belongs to tiny messages in reverse.
- **D, wrong:** the crossover is the whole point of the alpha-beta model.

Q4. A checkpoint holds a 70B BF16 model with Adam. Roughly how large is it, and what dominates?

- A. 140 GB, dominated by parameters
- B. 700 GB, dominated by optimizer state
- C. 700 GB, dominated by parameters
- D. 1.4 TB, dominated by gradient buffers

Answer: B. Parameters are 140 GB; Adam keeps two FP32 states per parameter (560 GB). Total ~700 GB, optimizer-dominated.

- **A, wrong:** forgets the optimizer, the classic checkpointing bug.
- **C, wrong:** right size, wrong component.
- **D, wrong:** gradients are not stored in checkpoints.

Q5. MTBF is 50 hours and a checkpoint save costs 90 seconds. What is the Daly-optimal save interval?

- A. 30 minutes
- B. 1.6 hours
- C. 5 hours
- D. 12 hours

Answer: B. $\sqrt{2 \times 90 \times 180,000} = \sqrt{32.4e6} = 5,692 \text{ s} = 1.58 \text{ h}$, about 1.6 hours.

- **A, wrong:** far too frequent; save overhead would dominate.
- **C, D, wrong:** too rare; each failure would waste too much work.

Q6. A 128-rank job shows a straggler factor of 1.4x. What is the most likely category of cause, and the first action?

- A. Recipe divergence; roll back the checkpoint
- B. One or a few slow nodes; identify and quarantine them
- C. Wrong parallelism layout; redesign TP/PP split
- D. Normal variance; raise the alert threshold

Answer: B. A 1.4x max/median ratio points at specific slow ranks: sick link, thermal throttling, or a dying GPU. Find them with per-rank step-time analysis and quarantine.

- **A, wrong:** loss behavior is not implicated by step-time skew.
- **C, wrong:** layout problems slow all ranks roughly equally.
- **D, wrong:** 1.4x is well above the 1.1x investigate threshold; normalizing it burns throughput.

Q7. In JAX, you annotate a weight matrix `PartitionSpec("data", "model")` on a mesh with axes `data = 2`, `model = 4`. How is the matrix split?

- A. Replicated on all 8 devices
- B. Rows split 2 ways, columns split 4 ways; each device holds one tile
- C. Rows split 4 ways, columns split 2 ways
- D. Split 8 ways along rows only

Answer: B. Each named axis shards the corresponding array axis: 2-way over rows, 4-way over columns, one tile per device.

- **A, wrong:** no sharding annotation means replication; a `PartitionSpec` always shards.
- **C, wrong:** swaps the axes.
- **D, wrong:** ignores the model axis.

Q8. A collective hang occurs. Logs show all ranks entered the same all-reduce, but one rank's NCCL log stops mid-operation while others time out waiting. What do you check first?

- A. The optimizer hyperparameters
- B. The fabric path to that rank: IB error counters, then bisect with `nccl-tests`
- C. The checkpoint format version
- D. The dataloader worker count

Answer: B. All ranks agreeing on the op rules out a code mismatch; one rank going silent mid-collective points at its network path or the GPU itself. Check IB counters, then isolate with `nccl-tests` on subsets.

- **A, wrong:** hyperparameters do not cause collective hangs.
- **C, wrong:** checkpoint format has nothing to do with in-flight collectives.
- **D, wrong:** dataloader stalls stop ranks before collectives, not mid-collective.

Q9. Why does tensor parallelism stay inside the NVLink domain?

- A. NVLink has lower latency for small messages
- B. A 10 GB per-layer all-reduce costs ~ 19 ms on NVLink vs ~ 350 ms on InfiniBand
- C. InfiniBand cannot do all-reduce
- D. Pipeline parallelism needs the IB links free

Answer: B. Ring traffic per rank is $2 \times 7/8 \times 10 = 17.5$ GB. At 900 GB/s that is 19 ms; at 50 GB/s it is 350 ms. A 2-second step absorbs 19 ms per layer but not 350 ms.

- **A, wrong:** the deciding factor is bandwidth on large per-layer messages, not latency.
- **C, wrong:** InfiniBand does all-reduce fine; it is just 18x slower per GPU.
- **D, wrong:** link sharing is a secondary concern, not the reason.

Q10. Async checkpointing cut your save stall from 200 s to 30 s. MTBF is 50.7 h. How does the Daly interval change?

- A. Unchanged; the interval depends only on MTBF
- B. From ~ 142 min to ~ 55 min
- C. From ~ 142 min to ~ 21 min
- D. It doubles, because faster saves are safer

Answer: B. $\sqrt{2 \times 200 \times 182,520} = 8,544 \text{ s} = 142 \text{ min}$. $\sqrt{2 \times 30 \times 182,520} = 3,309 \text{ s} = 55 \text{ min}$. Faster writes justify more frequent checkpoints.

- **A, wrong:** the interval balances write cost against failure cost; both matter.
- **C, wrong:** the interval scales with $\sqrt{C}$, not C : a 6.7x cost cut gives a 2.6x interval cut.
- **D, wrong:** backwards.

Q11. A run at 16,384 GPUs sees a failure every ~3 hours. What MTBF should a 1,000-GPU run on the same hardware plan for?

- A. ~3 hours
- B. ~12 hours
- C. ~50 hours
- D. ~800 hours

Answer: C. MTBF scales inversely with GPU count: $3 \text{ h} \times 16,384 / 1,000 = 49.2 \text{ h}$, about 50 h. (From the per-GPU-hour rate $1.97\text{e-}5$: $1/(1,000 \times 1.97\text{e-}5) = 50.7 \text{ h}$.)

- **A, wrong:** that is the 16k-GPU figure; fewer GPUs fail less often.
- **B, wrong:** undercounts the scaling.
- **D, wrong:** far too optimistic.

Q12. NCCL picks tree + LL for a 4 KB all-reduce across 64 ranks. Why is that correct?

- A. LL has higher bandwidth than Simple
- B. At 4 KB, latency dominates: tree finishes in ~61 us vs ~630 us for the ring
- C. Small messages cannot use the ring algorithm
- D. The tree uses fewer bytes on the wire

Answer: B. Alpha-beta math: the tree's $2 \times \log_2(64) = 12$ phases beat the ring's 126 steps when per-step latency dominates. Measured: 61 us vs 630 us.

- **A, wrong:** LL has LOWER bandwidth; it wins on latency, not bandwidth.
- **C, wrong:** the ring works fine on small messages; it is just slower.
- **D, wrong:** byte counts are similar; the difference is step count.

Q13. Your 70B run's checkpoint restore takes 40 minutes. The most likely design flaw is:

- A. The checkpoint interval is too long
- B. One rank reads the full 700 GB and redistributes it
- C. The optimizer state was not saved
- D. The filesystem is too fast

Answer: B. Sharded restore (each rank reads its own shard) finishes in about a minute at 8 ranks. A single-reader design serializes the whole 700 GB through one rank, then pays redistribution on top.

- **A, wrong:** interval affects write overhead and lost work, not restore time.
- **C, wrong:** missing optimizer state corrupts the resume; it does not slow the read.
- **D, wrong:** nonsense distractor.

Q14. Loss spikes to 6x its recent level for 2 steps, then recovers. The spike detector (median + 10xMAD, persist 5) stays silent. Correct behavior?

- A. No; any spike should trigger a rollback
- B. Yes; persist=5 filters single blips, and a 2-step excursion is noise
- C. No; the threshold k should be 2, not 10
- D. Yes, but only if the spike happened during warm-up

Answer: B. One or two bad steps are usually a noisy batch; the optimizer recovers on its own. Five sustained bad steps suggest poisoned momentum buffers. Rolling back on every blip would waste more than it saves.

- **A, wrong:** hair-trigger rollbacks cost hours; the persist counter exists to prevent them.
- **C, wrong:** $k = 10$ with MAD is the standard starting point; tune k on YOUR run's noise, but the persist logic is independent of k .
- **D, wrong:** warm-up is irrelevant to the persist logic.

Q15. ZeRO-3 moves about 1.5x the bytes of DDP per step. Why do large runs use it anyway?

- A. It reduces communication
- B. It shards parameters, grads, and optimizer state, so models that cannot fit under DDP become trainable
- C. It eliminates the all-reduce
- D. It works without NCCL

Answer: B. ZeRO-3 trades communication for memory. It moves $\sim 2.6\text{M}$ bytes per step vs $\sim 1.75\text{M}$ for DDP-class. Per-GPU memory drops to 16/8 bytes per param plus one live layer. For a 7B model that is 14 GB per GPU vs 38.5 GB under ZeRO-1. Fitting the model is the point.

- **A, wrong:** it moves MORE bytes, not fewer.
- **C, wrong:** it replaces one all-reduce with all-gather plus reduce-scatter.
- **D, wrong:** it still runs on NCCL.

Q16. A MoE layer routes 4096 tokens, top-2, to 64 experts with capacity factor 1.25. One expert receives 384 tokens. How many does it drop in training?

- A. 0
- B. 64
- C. 224
- D. 384

Answer: C. Expected load $4096 \times 2 / 64 = 128$ tokens. Buffer = $128 \times 1.25 = 160$ slots. Overflow = $384 - 160 = 224$ tokens dropped. The balancer nudges the router away from the hot expert on later steps.

- **A, wrong:** the buffer is finite; overflow drops.
- **B, wrong:** no basis in the arithmetic.
- **D, wrong:** the buffer still processes its 160.

Q17. Two all-reduce implementations move the same bytes, but one shows 27% exposed communication and the other 3%. The trace difference you expect:

- A. The 27% trace shows NCCL blocks interleaved with backward compute; the 3% shows one long block alone
- B. The 3% trace shows NCCL blocks interleaved with backward compute; the 27% shows one long block alone
- C. Exposed communication cannot be read from a trace
- D. Both traces look identical; exposure is a formula

Answer: B. Exposed communication is the fraction of comm time NOT hidden behind compute. Interleaved NCCL-under-backward blocks mean overlap; a lone all-reduce with empty compute rows means nothing hides it. Fix: smaller gradient buckets, larger microbatches, or earlier comm start.

- **A, wrong:** backwards.
- **C, wrong:** you measure exposure exactly on the trace.
- **D, wrong:** the overlap pattern is visible.

Exposed communication: what the trace shows

3% exposed: NCCL hides behind backward compute



blue: backward compute gold: NCCL, interleaved: nothing the compute rows did not already cover

27% exposed: one lone all-reduce, nothing hides it



blue: compute finishes, then the full all-reduce sits alone on the timeline
gold block = exposed: GPUs communicate while compute rows sit empty

The fix is overlap, not speed.

Smaller gradient buckets, larger microbatches, or an earlier comm start move gold blocks under blue ones. Interleaved NCCL under backward = healthy; a lone block with empty compute rows = exposed, and that is the number to kill.

What exposed communication looks like on a trace. Top: 3% exposed, NCCL blocks interleave under backward compute, so the network hides. Bottom: 27% exposed, one lone all-reduce sits with empty compute rows: that gold block is time no overlap can hide. The fix is overlap (smaller buckets, larger microbatches), not faster links.

Q18. On a TPU pod, why is "tensor parallelism inside the fast domain, data parallelism across" the wrong instinct?

- A. TPUs do not support tensor parallelism
- B. A pod is a switchless torus with uniform-ish ICI; there are no NVLink islands, so the cost model is hop count, not island vs uplink
- C. XLA forbids data parallelism
- D. TPU memory is too small for data parallelism

Answer: B. The DGX instinct assumes a fast 8-GPU island with slow uplinks. A torus gives every chip equal standing; blindly porting the island plan wastes the mesh. Name mesh axes to match the torus instead.

- **A, wrong:** TP maps to a mesh axis like any other strategy.
- **C, wrong:** data parallelism is a mesh axis too.
- **D, wrong:** memory size has nothing to do with the topology point.

Q19. After resuming from a checkpoint, the loss jumps discontinuously but then trains fine. The most likely resume bug is:

- A. The checkpoint was corrupt
- B. The LR scheduler was not restored to the checkpoint's step
- C. The model architecture changed
- D. The batch size changed

Answer: B. Schedules index by step. Resuming weights at step 5,000 with a scheduler at step 0 retrains with the wrong schedule, usually a higher LR than intended. That is the classic jump. The resume checklist requires the scheduler step to match.

- **A, wrong:** corruption usually breaks the restore or NaNs, not a clean jump.
- **C, D, wrong:** those change the setup visibly, not as a silent jump.

Q20. Context parallelism: Ulysses moves ~ 237 GB/step and ring attention moves $\sim 4x$ more bytes on a 128k-token setup, yet ring wins at extreme lengths. Why?

- A. Ring attention uses fewer FLOPs
- B. Overlapped bytes are free bytes, and ring keeps per-GPU memory flat while Ulysses has a transient all-tokens spike
- C. Ulysses is mathematically approximate
- D. Ring attention skips the softmax

Answer: B. Ring overlaps K/V block transfers with attention compute, and its memory per GPU stays flat as GPUs join. Ulysses briefly holds all tokens (for its head shard); at extreme lengths that spike OOMs. Byte counts lose to overlap plus memory scaling.

- **A, wrong:** FLOPs stay the same; attention math does not change.
- **C, wrong:** Ulysses is exact; the two all-to-alls just transpose the layout.
- **D, wrong:** ring computes the same softmax blockwise with running statistics.

Q21. A hang repro with 2 ranks and a tiny tensor passes, but the 1,000-GPU job hangs in the same collective. What does this rule out?

- A. Nothing; scale changes everything
- B. The launch and the basic code path: the bug is scale-dependent (config divergence across nodes, a sick node, or a fabric issue past some rank count)
- C. The network
- D. The collective algorithm

Answer: B. The 2-rank repro passing clears the launch and the code path. The standard next step: bisect the node set. The hang that only appears past some scale is usually env-var divergence on the added nodes or a sick node in the larger set.

- **A, wrong:** it rules out the cheap causes, which is the point of the repro.
- **C, wrong:** the fabric is still very much a suspect at scale.
- **D, wrong:** the algorithm does not change; the repro used it too.

Q22. FSDP all-gather traffic starts eating 25% of step time. Per the ladder, the next move is:

- A. More pipeline stages
- B. Add tensor parallelism inside the node, moving some traffic onto NVLink
- C. Larger global batch
- D. Switch to FP16

Answer: B. This is exactly the ladder's Q3->TP gate. When FSDP's all-gather traffic exceeds ~15% of step time, TP moves part of that traffic onto the fast intra-node fabric.

- **A, wrong:** PP is the last resort for very deep models, not the answer to comm fraction.
- **C, wrong:** a larger batch lengthens compute but does not fix the traffic mix.
- **D, wrong:** precision changes numerics, not the parallelism layout.

Q23. The NCCL tuner chose Ring + Simple for your TP all-reduces (tensors of ~200 KB). Step time regressed. What happened and what do you try?

- A. The tuner is always right; profile something else
- B. At 200 KB the tree + LL should win on latency; force NCCL_ALGO = Tree NCCL_PROTO = LL and measure
- C. Increase NCCL_MIN_NCHANNELS; more channels fix latency
- D. Switch to CollNet; in-network reduction is always faster

Answer: B. 200 KB is below the ~2.3 MB crossover: latency dominates and tree + LL should win. Forcing the override with a measurement in hand is one of the two legitimate override cases.

- **A, wrong:** the tuner is usually right, not always; the measurement decides.
- **C, wrong:** more channels split bandwidth; they do not fix per-step latency.
- **D, wrong:** CollNet needs SHARP-capable switches and validated config; it is not a default.

Q24. A GPU computes silently wrong results (no error raised) for an hour before anyone notices. Which safeguard would NOT have helped?

- A. Loss-spike detection
- B. More frequent checkpointing
- C. Per-layer activation max monitoring
- D. A second run with a different seed for comparison

Answer: B. Checkpointing preserves state faithfully, including corrupted state. It is the heartbeat, not the immune system. Detection (loss behavior, activation monitoring, redundancy) is what catches silent corruption.

- **A, wrong:** a spike detector can catch the loss signature of corruption.
- **C, wrong:** activation monitoring is exactly the kind of detection that helps.
- **D, wrong:** divergence between runs is a classic silent-corruption signal.

Q25. XLA recompiles your JAX training step on every batch. The likely cause and fix:

- A. The mesh is wrong; rebuild it
- B. Dynamic shapes: each new sequence length is a new (shape, dtype) compile key; pad to bucket boundaries
- C. The learning rate is too high
- D. gSPMD stays disabled

Answer: B. XLA compiles one binary per (shape, dtype). Varying lengths recompile every step. Bucketed padding (128, 256, 512) compiles once per bucket; the padding wastes a little compute, recompiling wastes far more.

- **A, wrong:** mesh shape affects sharding, not recompilation.
- **C, wrong:** LR affects training dynamics, not compilation.
- **D, wrong:** gSPMD is the partitioning pass; it does not control compile keys.

Q26. Your MFU dropped from 58% to 31% overnight with no config change. The trace shows one rank's backward pass stretching while its kernels look normal. First check:

- A. The loss scaler
- B. nvidia-smi clocks on that rank (thermal throttling), and whether the rank is slow inside collectives (NCCL timing) vs in kernels
- C. The checkpoint interval
- D. The router load balance

Answer: B. Normal-looking kernels with a stretching backward pass means the wait is upstream. Check throttling first: nvidia-smi clocks tell you in seconds. Then separate collective slowness from compute slowness with NCCL per-collective timing; slow inside collectives points at a fail-slow NIC. A rank that was fine for an hour then slows is usually heat, not code.

- **A, wrong:** the scaler affects numerics, not one rank's step time.
- **C, wrong:** checkpoint cadence does not stretch one rank's backward pass.
- **D, wrong:** MoE routing has nothing to do with a dense backward pass stretching.

Q27. DDP gradient buckets are tens of MB; TP per-layer all-reduces are KB to a few MB. What does NCCL pick for each, and why?

- A. Ring for both; the ring is always fastest
- B. Ring + Simple for DDP buckets, tree + LL for TP collectives: large messages are bandwidth-bound, small ones latency-bound
- C. Tree for both; fewer steps always wins
- D. It picks randomly per launch

Answer: B. This is the crossover applied: tens of MB sit above ~2.3 MB (ring territory), KB-to-MB sit below it (tree territory). The tuner follows the alpha-beta math.

- **A, wrong:** the ring loses by 10x on 4 KB messages.
- **C, wrong:** the tree loses by 6x on 1 GB messages.
- **D, wrong:** the tuner is deterministic given the inputs.

Q28. A 671B-parameter model in FP8 needs a full gradient sync. Roughly how long on a TPU pod with ~ 200 GB/s per chip of ICI?

- A. 0.07 s
- B. 0.7 s
- C. 6.7 s
- D. 67 s

Answer: C. Gradients are 671 GB; a ring all-reduce moves about $2 \times 671 \text{ GB} = 1.34 \text{ TB}$. At 200 GB/s that is $\sim 6.7 \text{ s}$. The lesson: at pod scale you never sync full parameters every step; you shard them.

- **A, wrong:** two orders of magnitude off; check the byte math.
- **B, wrong:** one order of magnitude off.
- **D, wrong:** one order of magnitude too high.

Q29. Pipeline bubble is 47% with 8 stages and 8 microbatches. You cannot add microbatches (activation memory). The honest fix is:

- A. Increase the microbatch size anyway; memory will sort itself out
- B. Fewer stages, more nodes per stage, or a different layout: the bubble formula $(p-1)/(m+p-1)$ is not negotiable
- C. Switch to a larger global batch; the bubble is batch-independent
- D. Overlap the bubble with checkpointing

Answer: B. The bubble is arithmetic, not a tuning problem. With m capped by memory, the only levers are fewer stages or a different strategy. Pretending otherwise rents GPUs to idle them.

- **A, wrong:** OOM is not a tuning outcome; it is a crash.
- **C, wrong:** the global batch does not appear in the bubble formula.
- **D, wrong:** checkpointing does not fill pipeline idle time.

Q30. After a failure, the run guardian restores the latest checkpoint but the loss jumps and stays high. Two candidate causes: (1) the checkpoint captured post-failure state, (2) the restore skipped the scheduler step. How do you tell them apart?

- A. You cannot; both look identical
- B. Check the loss-continuity against the pre-failure trend AND the scheduler step in the checkpoint metadata. A scheduler mismatch shows the wrong LR at the resume step, while a poisoned checkpoint shows a jump even with the right LR
- C. Restore an even older checkpoint; if it works, it was (1)
- D. Retrain from scratch

Answer: B. The resume checklist separates them. Verify the scheduler's step index against the checkpoint step: cause 2 shows a mismatch directly. Then compare the first post-resume steps to the pre-failure loss trend. Cause 1 (restoring poisoned state) jumps even when every index is correct. That is why the guardian keeps $K=3$ checkpoints and rolls back further on recurrence.

- **A, wrong:** the metadata check distinguishes them cleanly.
- **C, wrong:** restoring older is the treatment, not the diagnosis; check the metadata first.
- **D, wrong:** the most expensive possible response to a diagnosable problem.

Next: Section 5, the thirty numbers and rules on one page.

Section 5: One-page cheat sheet

The pack on one page: ratios, formulas, and order-of-magnitude thresholds.

Thirty numbers and rules. If you remember nothing else, remember these.

Efficiency

1. 6ND: training costs ~ 6 FLOPs per parameter per token (2 forward + 4 backward).
2. $MFU = 6 \times N \times \text{tok/s} / (n_{\text{gpus}} \times \text{peak})$. Always quote the peak you used.
3. H100 BF16 dense peak: 989 TFLOPS. A100: 312 TFLOPS.
4. Healthy MFU: 50-65% dense transformer, H100-class. Below 40%: profile now.
5. Goodput beats MFU: useful tokens after restarts, stalls, and stragglers.

Parallelism

1. Ladder: DDP -> FSDP -> TP -> PP. Memory fit, then comm fraction, then NVLink domain, then bubble.
2. Bubble = $(p-1)/(m+p-1)$. 8 stages x 8 micros = 47% idle. Keep micros well above stages.
3. TP stays inside the node. 10 GB all-reduce: 19 ms on NVLink, 350 ms on IB.
4. FSDP pays $2(N-1)/N$ twice per step (gather + scatter). At 512 ranks, 70B = ~ 559 GB per rank per step.
5. ZeRO-1/2/3 per-GPU bytes for 7B: 38.5 / 26.3 / 14.0 GB. ZeRO-3 comm ~ 1.5 x DDP.

Collectives

1. Ring all-reduce per rank: $2(N-1)/N \times M$ bytes ($\sim 2M$). All-to-all: $(N-1)/N \times M$ ($\sim M$).
2. $T = \text{steps} \times \alpha + \text{bytes}/\text{BW}$. Ring: $2(N-1)$ steps. Tree: $2 \times \log_2(N)$ phases.
3. Crossover ~ 2.3 MB (64 ranks, 40 GB/s, 5 us/step). Tree + LL below, ring + Simple above.
4. LL = 8 bytes/flag, CPU-pollled, microsecond latency. LL128 = 128 bytes/flag. Simple = bulk DMA.
5. Fabric cliff: NVLink 900 GB/s vs IB 50 GB/s = 18x. PCIe5 x16 ~ 64 GB/s for reference.
6. NCCL_DEBUG=INFO first. Flight recorder names the guilty rank. Shorten NCCL_TIMEOUT while debugging (default can be 30 min).

Checkpointing and recovery

1. 70B + Adam + BF16 = 700 GB (140 params + 560 optimizer). Budget the optimizer, not the model.
2. Daly: $\tau^* = \sqrt{2 \times C \times \mu}$. $C = 200$ s, $\mu = 50.7$ h -> every ~ 142 min. Async ($C = 30$ s) -> ~ 55 min.
3. MTBF scales inversely with GPU count: 1.97×10^{-5} per-GPU-hour. 1k GPUs -> ~ 50 h. 16k -> ~ 3 h.
4. Retention: last 3 + daily. Never one. A checkpoint never restored is a rumor.
5. Async: stage to pinned CPU (the only pause), one in-flight write, check the future.
6. Restore sharded: 8 ranks read in parallel ~ 1.3 min for 700 GB. One reader is not a plan.
7. Guardian loop under 10 min: detect, classify, isolate, resume. Quarantine nodes; they never rejoin silently.
8. Loss spike: median + $10 \times \text{MAD}$ for 5 steps -> stop, roll back pre-spike, skip bad batches, LR dip 2-4x briefly.

Failure and profiling

1. Llama 3 scale lesson: 16,384 GPUs, 54 days, 466 interruptions (~ 1 per 3 h), GPU faults 58.7%, effective time $> 90\%$ via automation.

2. Straggler factor = max/median rank step time. 1.1x investigate, 1.3x act. Never measure with barriers in the hot path.
 3. Fail-slow: > 1 h silent average, minutes to fix. Remediation economics favor fast detection.
 4. Hang order: your code (conditional collectives, uneven batches) -> slow rank -> fabric. Network is last, not first.
 5. Profile discipline: warm up tens of steps, distrust the mean, compare profiles to profiles, change one knob at a time.
 6. Precision: BF16 default (FP32 range). FP8 only with a validated recipe. MFU must use the matching peak.
-

Provenance

Last verified: September 2026.

- The MFU formula, 6ND rule, pipeline bubble, Daly interval, and alpha-beta model are standard results from the base volumes. All code outputs in this pack ran on this machine and match the printed outputs.
- Hardware numbers. H100 SXM BF16 dense ~989 TFLOPS and ~3,350 GB/s HBM. A100 80 GB 312 BF16 TFLOPS and 2,039 GB/s (vendor specs). NVLink 4 900 GB/s bidirectional, InfiniBand NDR400 400 Gb/s (50 GB/s per link), PCIe 5.0 x16 ~64 GB/s.
- NCCL algorithms (Ring, Tree, CollNet, NVLS), protocols (Simple, LL, LL128), and environment variables from NVIDIA's public NCCL documentation.
- Llama 3 405B run figures (16,384 GPUs, 54 days, 466 interruptions, 419 unexpected, GPU problems 58.7%, above 90% effective training time) from the public run report.
- Autoretry statistics (33.3% chain success, median 11-minute retry, 1.9 h vs 3.3 h downtime) from a public operational study of a ~500-GPU cluster.
- TPU specs (v6e ~918 BF16 TFLOPS, Ironwood 4,614 FP8 TFLOPS, pod sizes) from Google Cloud announcements and 2025-2026 coverage. JAX sharding API from current JAX documentation.
- **UNVERIFIED:** exact per-rank filesystem bandwidths, IB link latencies, and cloud-instance failure rates vary by cluster; the labs use representative values measured or estimated for illustration. The 200-second checkpoint stall in the Daly example is a representative estimate, not a measured site number. Fail-slow duration and remediation figures come from vendor and community reports, not peer-reviewed studies. DeepSeek-V3 deployment details appear via secondary write-ups of the paper.

Where to go next

The next pack in the sequence is [Crash: Evals and Safety](#). This pack trained the model. The next one answers the question this pack leaves open: how do you know the model is any good, and safe to ship? Cohen's kappa, Wilson intervals, pass@k, judge calibration, and red-teaming.

WATCH AND GO DEEPER

Two videos on the two hardest parts of distributed training: sharding everything with FSDP, and scaling expert-parallel MoE training to thousands of GPUs.

PyTorch FSDP: Architecture and Performance Optimization Strategies (Uplatz).

Opens with the memory wall that motivates everything: 112 GB of model state against 80 GB on an A100, then the DDP-to-FSDP contrast. Transcript verified, but honest note: it stays lighter on implementation detail than its description promises. Watch it for the intuition. Use the pack for the numbers.

Training MoEs at Scale with PyTorch (Databricks).

Brian Chu (Mosaic/Databricks) on training MoEs past a thousand GPUs. He covers expert parallelism, FSDP/HSDP combos, checkpointing pain, and the failure modes (hardware, S3, determinism) that only appear at scale. Transcript verified.

GO DEEPER

[ZeRO paper](#)

The three stages of optimizer-state, gradient, and parameter sharding, with the memory math behind the figure above.

[PyTorch FSDP tutorial](#)

The official hands-on: wrapping policies, mixed precision, and CPU offload in code.

[FSDP paper](#)

The systems paper: how FSDP shards and materializes parameters, and where the 1.5x comm cost comes from.